

Mathematics of Data: From Theory to Computation

Prof. Volkan Cevher
volkan.cevher@epfl.ch

Lecture 6: Unconstrained, smooth minimization III

Laboratory for Information and Inference Systems (LIONS)
École Polytechnique Fédérale de Lausanne (EPFL)

EE-556 (Fall 2018)



License Information for Mathematics of Data Slides

- ▶ This work is released under a [Creative Commons License](#) with the following terms:
- ▶ **Attribution**
 - ▶ The licensor permits others to copy, distribute, display, and perform the work. In return, licensees must give the original authors credit.
- ▶ **Non-Commercial**
 - ▶ The licensor permits others to copy, distribute, display, and perform the work. In return, licensees may not use the work for commercial purposes – unless they get the licensor's permission.
- ▶ **Share Alike**
 - ▶ The licensor permits others to distribute derivative works only under a license identical to the one that governs the licensor's work.
- ▶ [Full Text of the License](#)

Outline

- ▶ This lecture
 1. The quadratic case and conjugate gradient
 2. Other optimization methods
- ▶ Next lecture
 1. Stochastic gradient methods

Recommended reading

- ▶ Chapters 2, 3, 5, 6 in Nocedal, Jorge, and Wright, Stephen J., *Numerical Optimization*, Springer, 2006.
- ▶ Chapter 9 in Boyd, Stephen, and Vandenberghe, Lieven, *Convex optimization*, Cambridge university press, 2009.
- ▶ Chapter 1 in Bertsekas, Dimitris, *Nonlinear Programming*, Athena Scientific, 1999.
- ▶ Chapters 1, 2 and 4 in Nesterov, Yurii, *Introductory Lectures on Convex Optimization: A Basic Course*, Vol. 87, Springer, 2004.

Motivation

Motivation

This lecture covers some more advanced numerical methods for *unconstrained* and *smooth* convex minimization.

Recall: convex, unconstrained, smooth minimization

Problem (Mathematical formulation)

$$F^* := \min_{\mathbf{x} \in \mathbb{R}^p} \{F(\mathbf{x}) := f(\mathbf{x})\} \quad (1)$$

where f is *proper, closed, convex and twice differentiable*.

Note that (1) is unconstrained.

How do we design efficient optimization algorithms with accuracy-computation tradeoffs for this class of functions?

Linear systems

Problem (Solving a linear system)

Which is the best method for solving the linear system

$$\mathbf{Ax} = \mathbf{b} ?$$

Solving a linear system via optimization

To find a solution to the linear system, we can also solve the optimization problem

$$\min_{\mathbf{x}} f_{\mathbf{A},\mathbf{b}}(\mathbf{x}) := \frac{1}{2} \langle \mathbf{Ax}, \mathbf{x} \rangle - \langle \mathbf{b}, \mathbf{x} \rangle$$

which is seen to have a solution satisfying $\mathbf{Ax} = \mathbf{b}$ by solving $\nabla_{\mathbf{x}} f_{\mathbf{A},\mathbf{b}}(\mathbf{x}) = 0$.

- ▶ $f_{\mathbf{A},\mathbf{b}}$ is a quadratic function with **Lipschitz-gradient** ($L = \|\mathbf{A}\|$).
- ▶ If \mathbf{A} is a $p \times p$ symmetric positive definite matrix, (i.e., $\mathbf{A} = \mathbf{A}^T \succ 0$), $f_{\mathbf{A}}$ is also **strongly convex** ($\mu = \lambda_1(\mathbf{A})$, the smallest eigenvalue of \mathbf{A}).
- ▶ if \mathbf{A} is not symmetric, but full column rank, we can consider the system



$$\mathbf{A}^T \mathbf{Ax} = \mathbf{A}^T \mathbf{b}$$

which can be seen as: $\Phi \mathbf{x} = \mathbf{y}$ where Φ is symmetric and positive definite.

Linear systems

Remark

If Φ is diagonal and positive definite, given a starting point $\mathbf{x}^0 \in \text{dom}(f)$, successive minimization of $f_{\Phi, \mathbf{y}}(\mathbf{x})$ along the coordinate axes yield \mathbf{x}^* is at most p steps.



How can we adapt to the geometry of Φ ?

Conjugate gradients method - Φ symmetric and positive definite

Generate a set of *conjugate* directions $\{\mathbf{p}^0, \mathbf{p}^1, \dots, \mathbf{p}^{p-1}\}$ such that

$$\langle \mathbf{p}^i, \Phi \mathbf{p}^j \rangle = 0 \quad \text{for all } i \neq j \quad (\text{which also implies linear independence}).$$

Successively minimize $f_{\Phi, \mathbf{y}}$ along the individual conjugate directions. Let

$$\mathbf{r}^k = \Phi \mathbf{x}^k - \mathbf{y} \quad \text{and} \quad \mathbf{x}^{k+1} = \mathbf{x}^k + \alpha_k \mathbf{p}^k,$$

where α_k is the minimizer of $f_{\Phi, \mathbf{y}}(\mathbf{x})$ along $\mathbf{x}^k + \alpha \mathbf{p}^k$, i.e.,

$$\alpha_k = -\frac{\langle \mathbf{r}^k, \mathbf{p}^k \rangle}{\langle \mathbf{p}^k, \Phi \mathbf{p}^k \rangle}$$

Theorem

For any $\mathbf{x}^0 \in \mathbb{R}^p$ the sequence $\{\mathbf{x}^k\}$ generated by the conjugate directions algorithm converges to the solution \mathbf{x}^* of the linear system in **at most** p steps.

Intuition

The conjugate directions adapt to the geometry of the problem, taking the role of the canonical directions when Φ is a generic symmetric positive definite matrix.

Conjugate gradients method

Intuition

The conjugate directions adapt to the geometry of the problem, taking the role of the canonical directions when Φ is a generic symmetric positive definite matrix.

Back to diagonal

For a generic symmetric positive definite Φ , let us consider the variable $\bar{\mathbf{x}} := \mathbf{S}^{-1}\mathbf{x}$, with


$$\mathbf{S} = [\mathbf{p}^0, \dots, \mathbf{p}^{p-1}]$$

where $\{\mathbf{p}^k\}$ are the conjugate directions with respect to Φ . $f_{\Phi, \mathbf{y}}(\mathbf{x})$ now becomes

$$\bar{f}_{\Phi, \mathbf{y}}(\bar{\mathbf{x}}) := f_{\Phi, \mathbf{y}}(\mathbf{S}\bar{\mathbf{x}}) = \frac{1}{2} \langle \bar{\mathbf{x}}, (\mathbf{S}^T \Phi \mathbf{S}) \bar{\mathbf{x}} \rangle - \langle \mathbf{S}^T \mathbf{y}, \bar{\mathbf{x}} \rangle.$$

By the conjugacy property, $\langle \mathbf{p}^i, \Phi \mathbf{p}^j \rangle = 0$, $\forall i \neq j$, the matrix $\mathbf{S}^T \Phi \mathbf{S}$ is diagonal. Therefore, we can find the minimum of $\bar{f}(\bar{\mathbf{x}})$ in at most p steps along the canonical directions in $\bar{\mathbf{x}}$ space, which are the $\{\mathbf{p}^k\}$ directions in \mathbf{x} space.

Conjugate directions naturally adapt to the linear operator



Conjugate gradients method

Theorem

For any $\mathbf{x}^0 \in \mathbb{R}^p$ the sequence $\{\mathbf{x}^k\}$ generated by the conjugate directions algorithm converges to the solution \mathbf{x}^* of the linear system in **at most** p steps.

Proof.

Since $\{\mathbf{p}^k\}$ are linearly independent, they span \mathbb{R}^p . Therefore, we can write

$$\mathbf{x}^* - \mathbf{x}^0 = a_0 \mathbf{p}^0 + a_1 \mathbf{p}^1 + \cdots + a_{p-1} \mathbf{p}^{p-1}$$

for some values of the coefficients a_k . By multiplying on the left by $(\mathbf{p}^k)^T \Phi$ and using the conjugacy property, we obtain

$$a_k = \frac{\langle \mathbf{p}^k, \Phi(\mathbf{x}^* - \mathbf{x}^0) \rangle}{\langle \mathbf{p}^k, \Phi \mathbf{p}^k \rangle}.$$

Since $\mathbf{x}^k = \mathbf{x}^{k-1} + \alpha_{k-1} \mathbf{p}^{k-1}$, we have $\mathbf{x}^k = \mathbf{x}^0 + \alpha_0 \mathbf{p}^0 + \alpha_1 \mathbf{p}^1 + \cdots + \alpha_{k-1} \mathbf{p}^{k-1}$. By premultiplying by $(\mathbf{p}^k)^T \Phi$ and using the conjugacy property, we obtain $\langle \mathbf{p}^k, \Phi(\mathbf{x}^k - \mathbf{x}^0) \rangle = 0$ which implies

$$\langle \mathbf{p}^k, \Phi(\mathbf{x}^* - \mathbf{x}^0) \rangle = \langle \mathbf{p}^k, \Phi(\mathbf{x}^* - \mathbf{x}^k) \rangle = \langle \mathbf{p}^k, \mathbf{y} - \Phi \mathbf{x}^k \rangle = -\langle \mathbf{p}^k, \mathbf{r}^k \rangle$$

so that $a_k = -\frac{\langle \mathbf{p}^k, \mathbf{r}^k \rangle}{\langle \mathbf{p}^k, \Phi \mathbf{p}^k \rangle} = \alpha_k$. □

Conjugate gradients method

How can we efficiently generate a set of conjugate directions?

Iteratively generate the new descent direction \mathbf{p}^k from the previous one:

$$\mathbf{p}^k = -\mathbf{r}^k + \beta_k \mathbf{p}^{k-1}$$

For ensuring conjugacy $\langle \mathbf{p}^k, \Phi \mathbf{p}^{k-1} \rangle = 0$, we need to choose β_k as

$$\beta_k = \frac{\langle \mathbf{r}^k, \Phi \mathbf{p}^{k-1} \rangle}{\langle \mathbf{p}^{k-1}, \Phi \mathbf{p}^{k-1} \rangle} .$$

Lemma

The directions $\{\mathbf{p}^0, \mathbf{p}^1, \dots, \mathbf{p}^p\}$ form a conjugate directions set.

Conjugate gradients method

Conjugate gradients (CG) method

1 Initialization:

1.a Choose $\mathbf{x}^0 \in \text{dom}(f)$ arbitrarily as a starting point.

1.b Set $\mathbf{r}^0 = \Phi \mathbf{x}^0 - \mathbf{y}$, $\mathbf{p}^0 = -\mathbf{r}^0$, $k = 0$.

2. While $\mathbf{r}^k \neq \mathbf{0}$, generate a sequence $\{\mathbf{x}^k\}_{k \geq 0}$ as:

$$\begin{aligned}\alpha_k &= -\frac{\langle \mathbf{r}^k, \mathbf{p}^k \rangle}{\langle \mathbf{p}^k, \Phi \mathbf{p}^k \rangle} \\ \mathbf{x}^{k+1} &= \mathbf{x}^k + \alpha_k \mathbf{p}^k \\ \mathbf{r}^{k+1} &= \Phi \mathbf{x}^{k+1} - \mathbf{y} \\ \beta_{k+1} &= \frac{\langle \mathbf{r}^{k+1}, \Phi \mathbf{p}^k \rangle}{\langle \mathbf{p}^k, \Phi \mathbf{p}^k \rangle} \\ \mathbf{p}^{k+1} &= -\mathbf{r}^{k+1} + \beta_{k+1} \mathbf{p}^k \\ k &= k + 1\end{aligned}$$

Theorem

Since the directions $\{\mathbf{p}^0, \mathbf{p}^1, \dots, \mathbf{p}^k\}$ are conjugate, CG converges in at most p steps.

Other properties of the conjugate gradient method

Theorem

If Φ has only r distinct eigenvalues, then the CG iterations will terminate at the solution in at most r iterations.

Theorem

If Φ has eigenvalues $\lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_p$, we have that

$$\|\mathbf{x}^{k+1} - \mathbf{x}^*\|_{\Phi} \leq \left(\frac{\lambda_{p-k} - \lambda_1}{\lambda_{p-k} + \lambda_1} \right) \|\mathbf{x}^0 - \mathbf{x}^*\|_{\Phi},$$

where the local norm is given by $\|\mathbf{x}\|_{\Phi} = \sqrt{\mathbf{x}^T \Phi \mathbf{x}}$.

Theorem

Conjugate gradients algorithm satisfy the following iteration invariant for the solution of $\Phi \mathbf{x} = \mathbf{y}$

$$\|\mathbf{x}^{k+1} - \mathbf{x}^*\|_{\Phi} \leq 2 \left(\frac{\sqrt{\kappa(\Phi)} - 1}{\sqrt{\kappa(\Phi)} + 1} \right)^k \|\mathbf{x}^0 - \mathbf{x}^*\|_{\Phi},$$

where the condition number of Φ is defined as $\kappa(\Phi) := \|\Phi\| \|\Phi^{-1}\| = \frac{\lambda_p}{\lambda_1}$.

GD and AGD for the quadratic case: choice of the step size

Gradient Descent

$$\alpha_k = \frac{2}{L + \mu} \quad \text{with } L = \lambda_p(\Phi) \text{ and } \mu = \lambda_1(\Phi)$$

Steepest descent

Choose α_k so as to minimize $f(\mathbf{x}^{k+1})$.

$$\alpha_k = \frac{\|\nabla f(\mathbf{x}^k)\|^2}{\langle \nabla f(\mathbf{x}^k), \Phi \nabla f(\mathbf{x}^k) \rangle} \quad (1)$$

Barzilai-Borwein

$$\alpha_k = \frac{\|\nabla f(\mathbf{x}^{k-1})\|^2}{\langle \nabla f(\mathbf{x}^{k-1}), \Phi \nabla f(\mathbf{x}^{k-1}) \rangle} \quad (2)$$

The quadratic case - convergence rates summary

Convergence rates

$$\text{Gradient descent } \left(\alpha_k = \frac{2}{L+\mu} \right) : \quad \| \mathbf{x}^k - \mathbf{x}^* \|_2 \leq \left(\frac{\lambda_p - \lambda_1}{\lambda_p} \right)^k \| \mathbf{x}^0 - \mathbf{x}^* \|_2$$

$$\text{Steepest descent:} \quad \| \mathbf{x}^{k+1} - \mathbf{x}^* \|_{\Phi} \leq \left(\frac{\lambda_p - \lambda_1}{\lambda_p + \lambda_1} \right)^k \| \mathbf{x}^0 - \mathbf{x}^* \|_{\Phi}$$



$$\text{Barzilai-Borwein } (\lambda_p < 2\lambda_1) : \quad \| \mathbf{x}^{k+1} - \mathbf{x}^* \|_2 \leq \left(\frac{\lambda_p - \lambda_1}{\lambda_1} \right)^k \| \mathbf{x}^0 - \mathbf{x}^* \|_2$$

$$\text{AGD-}\mu\text{L:} \quad \| \mathbf{x}^k - \mathbf{x}^* \|_2 \leq \left(\frac{\sqrt{\lambda_p} - \sqrt{\lambda_1}}{\sqrt{\lambda_p}} \right)^{\frac{k}{2}} \| \mathbf{x}^0 - \mathbf{x}^* \|_2$$



$$\text{Conjugate gradient method:} \quad \| \mathbf{x}^{k+1} - \mathbf{x}^* \|_{\Phi} \leq \left(\frac{\sqrt{\lambda_p} - \sqrt{\lambda_1}}{\sqrt{\lambda_p} + \sqrt{\lambda_1}} \right)^k \| \mathbf{x}^0 - \mathbf{x}^* \|_{\Phi}$$

Example: Quadratic function


Case 1: $n = p = 1000, \kappa(\mathbf{A}) = 100$




Case 2: $n = p = 1000, \kappa(\mathbf{A}) = 1000$




How can we better adapt to the local geometry?



How can we better adapt to the local geometry?



How can we better adapt to the local geometry?



Variable metric gradient descent algorithm

Variable metric gradient descent algorithm

1. Choose $\mathbf{x}^0 \in \mathbb{R}^p$ as a starting point and $\mathbf{H}_0 \succ 0$.
2. For $k = 0, 1, \dots$, perform:

$$\begin{cases} \mathbf{d}^k & := -\mathbf{H}_k^{-1} \nabla f(\mathbf{x}^k), \\ \mathbf{x}^{k+1} & := \mathbf{x}^k + \alpha_k \mathbf{d}^k, \end{cases}$$

where $\alpha_k \in (0, 1]$ is a given step size.

3. Update $\mathbf{H}_{k+1} \succ 0$ if necessary.

Variable metric gradient descent algorithm

Variable metric gradient descent algorithm

1. Choose $\mathbf{x}^0 \in \mathbb{R}^p$ as a starting point and $\mathbf{H}_0 \succ 0$.
2. For $k = 0, 1, \dots$, perform:

$$\begin{cases} \mathbf{d}^k &:= -\mathbf{H}_k^{-1} \nabla f(\mathbf{x}^k), \\ \mathbf{x}^{k+1} &:= \mathbf{x}^k + \alpha_k \mathbf{d}^k, \end{cases}$$

where $\alpha_k \in (0, 1]$ is a given step size.

3. Update $\mathbf{H}_{k+1} \succ 0$ if necessary.

Common choices of the variable metric \mathbf{H}_k

- ▶ $\mathbf{H}_k := \lambda_k \mathbf{I} \implies$ gradient descent method.
- ▶ $\mathbf{H}_k := \mathbf{D}_k$ (a positive diagonal matrix) \implies scaled gradient descent method.
- ▶ $\mathbf{H}_k := \nabla^2 f(\mathbf{x}^k) \implies$ Newton method.
- ▶ $\mathbf{H}_k \approx \nabla^2 f(\mathbf{x}^k) \implies$ quasi-Newton method.

Newton method

- ▶ **Fast** (local) convergence but **expensive** per iteration cost
- ▶ **Useful** when **warm-started** near a solution

Newton method

- ▶ **Fast** (local) convergence but **expensive** per iteration cost
- ▶ **Useful** when **warm-started** near a solution

Local quadratic approximation using the Hessian

- ▶ Obtain a local quadratic approximation using the second-order Taylor series approximation to $f(\mathbf{x}^k + \mathbf{p})$:

$$f(\mathbf{x}^k + \mathbf{p}) \approx f(\mathbf{x}^k) + \langle \mathbf{p}, \nabla f(\mathbf{x}^k) \rangle + \frac{1}{2} \langle \mathbf{p}, \nabla^2 f(\mathbf{x}^k) \mathbf{p} \rangle$$

Newton method

- ▶ **Fast** (local) convergence but **expensive** per iteration cost
- ▶ **Useful** when **warm-started** near a solution

Local quadratic approximation using the Hessian

- ▶ Obtain a local quadratic approximation using the second-order Taylor series approximation to $f(\mathbf{x}^k + \mathbf{p})$:

$$f(\mathbf{x}^k + \mathbf{p}) \approx f(\mathbf{x}^k) + \langle \mathbf{p}, \nabla f(\mathbf{x}^k) \rangle + \frac{1}{2} \langle \mathbf{p}, \nabla^2 f(\mathbf{x}^k) \mathbf{p} \rangle$$

- ▶ The Newton direction is the vector \mathbf{p}^k that minimizes $f(\mathbf{x}^k + \mathbf{p})$; assuming the Hessian $\nabla^2 f_k$ to be **positive definite**, :

$$\nabla^2 f(\mathbf{x}^k) \mathbf{p}^k = -\nabla f(\mathbf{x}^k) \quad \Leftrightarrow \quad \mathbf{p}^k = -\left(\nabla^2 f(\mathbf{x}^k)\right)^{-1} \nabla f(\mathbf{x}^k)$$

Newton method

- ▶ **Fast** (local) convergence but **expensive** per iteration cost
- ▶ **Useful** when **warm-started** near a solution

Local quadratic approximation using the Hessian

- ▶ Obtain a local quadratic approximation using the second-order Taylor series approximation to $f(\mathbf{x}^k + \mathbf{p})$:

$$f(\mathbf{x}^k + \mathbf{p}) \approx f(\mathbf{x}^k) + \langle \mathbf{p}, \nabla f(\mathbf{x}^k) \rangle + \frac{1}{2} \langle \mathbf{p}, \nabla^2 f(\mathbf{x}^k) \mathbf{p} \rangle$$

- ▶ The Newton direction is the vector \mathbf{p}^k that minimizes $f(\mathbf{x}^k + \mathbf{p})$; assuming the Hessian $\nabla^2 f_k$ to be **positive definite**, :

$$\nabla^2 f(\mathbf{x}^k) \mathbf{p}^k = -\nabla f(\mathbf{x}^k) \quad \Leftrightarrow \quad \mathbf{p}^k = -\left(\nabla^2 f(\mathbf{x}^k)\right)^{-1} \nabla f(\mathbf{x}^k)$$

- ▶ A unit step-size $\alpha_k = 1$ can be chosen near convergence:

$$\mathbf{x}^{k+1} = \mathbf{x}^k - \left(\nabla^2 f(\mathbf{x}^k)\right)^{-1} \nabla f(\mathbf{x}^k) .$$

Newton method

- ▶ **Fast** (local) convergence but **expensive** per iteration cost
- ▶ **Useful** when **warm-started** near a solution

Local quadratic approximation using the Hessian

- ▶ Obtain a local quadratic approximation using the second-order Taylor series approximation to $f(\mathbf{x}^k + \mathbf{p})$:

$$f(\mathbf{x}^k + \mathbf{p}) \approx f(\mathbf{x}^k) + \langle \mathbf{p}, \nabla f(\mathbf{x}^k) \rangle + \frac{1}{2} \langle \mathbf{p}, \nabla^2 f(\mathbf{x}^k) \mathbf{p} \rangle$$

- ▶ The Newton direction is the vector \mathbf{p}^k that minimizes $f(\mathbf{x}^k + \mathbf{p})$; assuming the Hessian $\nabla^2 f_k$ to be **positive definite**, :

$$\nabla^2 f(\mathbf{x}^k) \mathbf{p}^k = -\nabla f(\mathbf{x}^k) \quad \Leftrightarrow \quad \mathbf{p}^k = -\left(\nabla^2 f(\mathbf{x}^k)\right)^{-1} \nabla f(\mathbf{x}^k)$$

- ▶ A unit step-size $\alpha_k = 1$ can be chosen near convergence:

$$\mathbf{x}^{k+1} = \mathbf{x}^k - \left(\nabla^2 f(\mathbf{x}^k)\right)^{-1} \nabla f(\mathbf{x}^k) .$$

Remark

- ▶ For $f \in \mathcal{F}_{L,1}^{2,1}$ but $f \notin \mathcal{F}_{L,\mu}^{2,1}$, the Hessian may not always be positive definite.

(Local) Convergence of Newton method

Lemma

Assume f is a twice differentiable convex function with minimum at \mathbf{x}^* such that:

- ▶ $\nabla^2 f(\mathbf{x}^*) \succeq \mu \mathbf{I}$ for some $\mu > 0$,
- ▶ $\|\nabla^2 f(\mathbf{x}) - \nabla^2 f(\mathbf{y})\|_{2 \rightarrow 2} \leq M \|\mathbf{x} - \mathbf{y}\|_2$ for some constant $M > 0$ and all $\mathbf{x}, \mathbf{y} \in \text{dom}(f)$.

Moreover, assume the starting point $\mathbf{x}^0 \in \text{dom}(f)$ is such that $\|\mathbf{x}^0 - \mathbf{x}^*\|_2 < \frac{2\mu}{3M}$. Then, the Newton method iterates converge **quadratically**:

$$\|\mathbf{x}^{k+1} - \mathbf{x}^*\| \leq \frac{M \|\mathbf{x}^k - \mathbf{x}^*\|_2^2}{2 \left(\mu - M \|\mathbf{x}^k - \mathbf{x}^*\|_2 \right)}.$$

Remark

This is the fastest convergence rate we have seen so far, but it requires to solve a $p \times p$ linear system at each iteration, $\nabla^2 f(\mathbf{x}^k) \mathbf{p}^k = -\nabla f(\mathbf{x}^k)$!

Locally quadratic convergence of the Newton method-I

Newton's method local quadratic convergence - Proof [2]

Since $\nabla f(\mathbf{x}^*) = 0$ we have

$$\begin{aligned}\mathbf{x}^{k+1} - \mathbf{x}^* &= \mathbf{x}^k - \mathbf{x}^* - (\nabla^2 f(\mathbf{x}^k))^{-1} \nabla f(\mathbf{x}^k) \\ &= (\nabla^2 f(\mathbf{x}^k))^{-1} \left(\nabla^2 f(\mathbf{x}^k)(\mathbf{x}^k - \mathbf{x}^*) - (\nabla f(\mathbf{x}^k) - \nabla f(\mathbf{x}^*)) \right)\end{aligned}$$

By Taylor's theorem, we also have

$$\nabla f(\mathbf{x}^k) - \nabla f(\mathbf{x}^*) = \int_0^1 \nabla^2 f(\mathbf{x}^k + t(\mathbf{x}^* - \mathbf{x}^k))(\mathbf{x}^k - \mathbf{x}^*) dt$$

Combining the two above, we obtain

$$\begin{aligned}& \|\nabla^2 f(\mathbf{x}^k)(\mathbf{x}^k - \mathbf{x}^*) - (\nabla f(\mathbf{x}^k) - \nabla f(\mathbf{x}^*))\| \\ &= \left\| \int_0^1 \left(\nabla^2 f(\mathbf{x}^k) - \nabla^2 f(\mathbf{x}^k + t(\mathbf{x}^* - \mathbf{x}^k)) \right) (\mathbf{x}^k - \mathbf{x}^*) dt \right\| \\ &\leq \int_0^1 \left\| \nabla^2 f(\mathbf{x}^k) - \nabla^2 f(\mathbf{x}^k + t(\mathbf{x}^* - \mathbf{x}^k)) \right\| \|\mathbf{x}^k - \mathbf{x}^*\| dt \\ &\leq M \|\mathbf{x}^k - \mathbf{x}^*\|^2 \int_0^1 t dt = \frac{1}{2} M \|\mathbf{x}^k - \mathbf{x}^*\|^2\end{aligned}$$

Locally quadratic convergence of the Newton method–II

Newton's method local quadratic convergence - Proof [2].

- ▶ Recall

$$\mathbf{x}^{k+1} - \mathbf{x}^* = (\nabla^2 f(\mathbf{x}^k))^{-1} \left(\nabla^2 f(\mathbf{x}^k)(\mathbf{x}^k - \mathbf{x}^*) - (\nabla f(\mathbf{x}^k) - \nabla f(\mathbf{x}^*)) \right)$$

$$\|\nabla^2 f(\mathbf{x}^k)(\mathbf{x}^k - \mathbf{x}^*) - (\nabla f(\mathbf{x}^k) - \nabla f(\mathbf{x}^*))\| \leq \frac{1}{2}M\|\mathbf{x}^k - \mathbf{x}^*\|^2$$

- ▶ Since $\nabla^2 f(\mathbf{x}^*)$ is nonsingular, there must exist a radius r such that $\|(\nabla^2 f(\mathbf{x}^k))^{-1}\| \leq 2\|(\nabla^2 f(\mathbf{x}^*))^{-1}\|$ for all \mathbf{x}^k with $\|\mathbf{x}^k - \mathbf{x}^*\| \leq r$.
- ▶ Substituting, we obtain

$$\|\mathbf{x}^{k+1} - \mathbf{x}^*\| \leq M\|(\nabla^2 f(\mathbf{x}^*))^{-1}\|\|\mathbf{x}^k - \mathbf{x}^*\|^2 = \tilde{M}\|\mathbf{x}^k - \mathbf{x}^*\|^2,$$

where $\tilde{M} = M\|(\nabla^2 f(\mathbf{x}^*))^{-1}\|$.

- ▶ If we choose $\|\mathbf{x}^0 - \mathbf{x}^*\| \leq \min(r, 1/(2\tilde{M}))$, we obtain by induction that the iterates \mathbf{x}^k converge quadratically to \mathbf{x}^* .



Example: Logistic regression

Problem (Logistic regression)



Given $\mathbf{A} \in \{0, 1\}^{n \times p}$ and $\mathbf{b} \in \{-1, +1\}^n$, solve:

$$f^* := \min_{\mathbf{x}, \beta} \left\{ f(\mathbf{x}) := \frac{1}{n} \sum_{j=1}^n \log \left(1 + \exp \left(-\mathbf{b}_j (\mathbf{a}_j^T \mathbf{x} + \beta) \right) \right) \right\}.$$

Real data

- ▶ Real data: w5a with $n = 9888$ data points, $p = 300$ features
- ▶ Available at <http://www.csie.ntu.edu.tw/~cjlin/libsvmtools/datasets/binary.html>.

Example: Logistic regression - numerical results



Parameters

- ▶ Newton's method: maximum number of iterations 200, tolerance 10^{-6} .
- ▶ For accelerated gradient method: maximum number of iterations 20000, tolerance 10^{-6} .
- ▶ Ground truth: Get a high accuracy approximation of \mathbf{x}^* and f^* by applying Newton's method for 200 iterations.

Quasi-Newton methods

Quasi-Newton methods use an approximate Hessian oracle and can be more scalable.

- ▶ Useful for $f(\mathbf{x}) := \sum_{i=1}^n f_i(\mathbf{x})$ with $n \gg p$.

Main ingredients

Quasi-Newton direction:

$$\mathbf{p}^k = -\mathbf{H}_k^{-1} \nabla f(\mathbf{x}^k) = -\mathbf{B}_k \nabla f(\mathbf{x}^k).$$

- ▶ Matrix \mathbf{H}_k , or its inverse \mathbf{B}_k , undergoes low-rank updates:
 - ▶ Rank 1 or 2 updates: famous Broyden–Fletcher–Goldfarb–Shanno (BFGS) algorithm.
 - ▶ Limited memory BFGS (L-BFGS).
- ▶ Line-search: The step-size α_k is chosen to satisfy the **Wolfe conditions**:

$$f(\mathbf{x}^k + \alpha_k \mathbf{p}^k) \leq f(\mathbf{x}^k) + c_1 \alpha_k \langle \nabla f(\mathbf{x}^k), \mathbf{p}^k \rangle \quad (\text{sufficient decrease})$$

$$\langle \nabla f(\mathbf{x}^k + \alpha_k \mathbf{p}^k), \mathbf{p}^k \rangle \geq c_2 \langle \nabla f(\mathbf{x}^k), \mathbf{p}^k \rangle \quad (\text{curvature condition})$$

with $0 < c_1 < c_2 < 1$. For quasi-Newton methods, we usually use $c_1 = 0.1$.

- ▶ Convergence is guaranteed under the Dennis & Moré condition [1].
- ▶ For more details on quasi-Newton methods, see Nocedal&Wright's book [2].

*Quasi-Newton methods

How do we update \mathbf{B}_{k+1} ?

Suppose we have (note the coordinate change from \mathbf{p} to $\bar{\mathbf{p}}$)

$$m_{k+1}(\bar{\mathbf{p}}) := f(\mathbf{x}^{k+1}) + \langle \nabla f(\mathbf{x}^{k+1}), \bar{\mathbf{p}} - \mathbf{x}^{k+1} \rangle + \frac{1}{2} \langle \mathbf{B}_{k+1}(\bar{\mathbf{p}} - \mathbf{x}^{k+1}), (\bar{\mathbf{p}} - \mathbf{x}^{k+1}) \rangle.$$

We require the gradient of m_{k+1} to match the gradient of f at \mathbf{x}^k and \mathbf{x}^{k+1} .

- ▶ $\nabla m_{k+1}(\mathbf{x}^{k+1}) = \nabla f(\mathbf{x}^{k+1})$ as desired;
- ▶ For \mathbf{x}^k , we have

$$\nabla m_{k+1}(\mathbf{x}^k) = \nabla f(\mathbf{x}^{k+1}) + \mathbf{B}_{k+1}(\mathbf{x}^k - \mathbf{x}^{k+1})$$

which must be equal to $\nabla f(\mathbf{x}^k)$.

- ▶ Rearranging, we have that \mathbf{B}_{k+1} must satisfy the **secant equation**

$$\mathbf{B}_{k+1} \mathbf{s}^k = \mathbf{y}^k$$

where $\mathbf{s}^k = \mathbf{x}^{k+1} - \mathbf{x}^k$ and $\mathbf{y}^k = \nabla f(\mathbf{x}^{k+1}) - \nabla f(\mathbf{x}^k)$.

- ▶ The secant equation can be satisfied with a positive definite matrix \mathbf{B}_{k+1} only if $\langle \mathbf{s}^k, \mathbf{y}^k \rangle > 0$, which is guaranteed to hold if the step-size α_k satisfies the Wolfe conditions.

*Quasi-Newton methods

BFGS method [2] (from Broyden, Fletcher, Goldfarb & Shanno)

The BFGS method arises from directly updating $\mathbf{H}_k = \mathbf{B}_k^{-1}$. The update on the inverse \mathbf{B} is found by solving

$$\min_{\mathbf{H}} \|\mathbf{H} - \mathbf{H}_k\|_{\mathbf{W}} \quad \text{subject to } \mathbf{H} = \mathbf{H}^T \text{ and } \mathbf{H}\mathbf{y}^k = \mathbf{s}^k \quad (3)$$

The solution is a rank-2 update of the matrix \mathbf{H}_k :

$$\mathbf{H}_{k+1} = \mathbf{V}_k^T \mathbf{H}_k \mathbf{V}_k + \eta_k \mathbf{s}^k (\mathbf{s}^k)^T,$$

where $\mathbf{V}_k = \mathbf{I} - \eta_k \mathbf{y}^k (\mathbf{s}^k)^T$.

- Initialization of \mathbf{H}_0 is an art. We can choose to set it to be an approximation of $\nabla^2 f(\mathbf{x}^0)$ obtained by finite differences or just a multiple of the identity matrix.

*Quasi-Newton methods

BFGS method [2] (from Broyden, Fletcher, Goldfarb & Shanno)

The BFGS method arises from directly updating $\mathbf{H}_k = \mathbf{B}_k^{-1}$. The update on the inverse \mathbf{B} is found by solving

$$\min_{\mathbf{H}} \|\mathbf{H} - \mathbf{H}_k\|_{\mathbf{W}} \quad \text{subject to } \mathbf{H} = \mathbf{H}^T \text{ and } \mathbf{H}\mathbf{y}^k = \mathbf{s}^k \quad (3)$$

The solution is a rank-2 update of the matrix \mathbf{H}_k :

$$\mathbf{H}_{k+1} = \mathbf{V}_k^T \mathbf{H}_k \mathbf{V}_k + \eta_k \mathbf{s}^k (\mathbf{s}^k)^T,$$

where $\mathbf{V}_k = \mathbf{I} - \eta_k \mathbf{y}^k (\mathbf{s}^k)^T$.

Theorem (Convergence of BFGS)

Let $f \in \mathcal{C}^2$. Assume that the BFGS sequence $\{\mathbf{x}^k\}$ converges to a point \mathbf{x}^* and $\sum_{k=1}^{\infty} \|\mathbf{x}^k - \mathbf{x}^*\| \leq \infty$. Assume also that $\nabla^2 f(\mathbf{x})$ is Lipschitz continuous at \mathbf{x}^* . Then \mathbf{x}^k converges to \mathbf{x}^* at a **superlinear** rate.

Remarks

The proof shows that given the assumptions, the BFGS updates for \mathbf{B}_k satisfy the Dennis & Moré condition, which in turn implies superlinear convergence.

L-BFGS

Challenges for BFGS

- ▶ BFGS approach stores and applies a dense $p \times p$ matrix \mathbf{H}_k .
- ▶ When p is very large, \mathbf{H}_k can prohibitively expensive to store and apply.

L(imited memory)-BFGS

- ▶ Do not store \mathbf{H}_k , but keep only the m most recent pairs $\{(\mathbf{s}^i, \mathbf{y}^i)\}$.
- ▶ Compute $\mathbf{H}_k \nabla f(\mathbf{x}_k)$ by performing a sequence of operations with \mathbf{s}^i and \mathbf{y}^i :
 - ▶ Choose a temporary initial approximation \mathbf{H}_k^0 .
 - ▶ Recursively apply $\mathbf{H}_{k+1} = \mathbf{V}_k^T \mathbf{H}_k \mathbf{V}_k + \eta_k \mathbf{s}^k (\mathbf{s}^k)^T$, m times starting from \mathbf{H}_k^0 :

$$\begin{aligned} \mathbf{H}_k &= \left(\mathbf{V}_{k-1}^T \cdots \mathbf{V}_{k-m}^T \right) \mathbf{H}_k^0 \left(\mathbf{V}_{k-m} \cdots \mathbf{V}_{k-1} \right) \\ &\quad + \eta_{k-m} \left(\mathbf{V}_{k-1}^T \cdots \mathbf{V}_{k-m+1}^T \right) \mathbf{s}^{k-m} (\mathbf{s}^{k-m})^T \left(\mathbf{V}_{k-m+1} \cdots \mathbf{V}_{k-1} \right) \\ &\quad + \cdots \\ &\quad + \eta_{k-1} \mathbf{s}^{k-1} (\mathbf{s}^{k-1})^T \end{aligned}$$

- ▶ From the previous expression, we can compute $\mathbf{H}_k \nabla f(\mathbf{x}^k)$ recursively.
- ▶ Replace the oldest element in $\{\mathbf{s}^i, \mathbf{y}^i\}$ with $(\mathbf{s}^k, \mathbf{y}^k)$.
- ▶ From practical experience, $m \in (3, 50)$ does the trick.

L-BFGS: A quasi-Newton method

Procedure for computing $\mathbf{H}_k \nabla f(\mathbf{x}^k)$	
0.	Recall $\eta_k = 1/\langle \mathbf{y}^k, \mathbf{s}^k \rangle$.
1.	$\mathbf{q} = \nabla f(\mathbf{x}^k)$.
2.	For $i = k - 1, \dots, k - m$ $\alpha_i = \eta_i \langle \mathbf{s}^i, \mathbf{q} \rangle$ $\mathbf{q} = \mathbf{q} - \alpha_i \mathbf{y}^i.$
3.	$\mathbf{r} = \mathbf{H}_k^0 \mathbf{q}$.
4.	For $i = k - m, \dots, k - 1$ $\beta = \eta_i \langle \mathbf{y}^i, \mathbf{r} \rangle$ $\mathbf{r} = \mathbf{r} + (\alpha_i - \beta) \mathbf{s}^i.$
5.	$\mathbf{H}_k \nabla f(\mathbf{x}^k) = \mathbf{r}$.

Remarks

- ▶ Apart from the step $\mathbf{r} = \mathbf{H}_k^0 \mathbf{q}$, the algorithm requires only $4mp$ multiplications.
- ▶ If \mathbf{H}_k^0 is chosen to be diagonal, another p multiplications are needed.
- ▶ An effective initial choice is $\mathbf{H}_k^0 = \gamma_k \mathbf{I}$, where

$$\gamma_k = \frac{\langle \mathbf{s}^{k-1}, \mathbf{y}^{k-1} \rangle}{\langle \mathbf{y}^{k-1}, \mathbf{y}^{k-1} \rangle}$$

L-BFGS: A quasi-Newton method



L-BFGS

1. Choose starting point \mathbf{x}^0 and $m > 0$.
2. For $k = 0, 1, \dots$
 - 2.a Choose \mathbf{H}_k^0 .
 - 2.b Compute $\mathbf{p}^k = -\mathbf{H}_k \nabla f(\mathbf{x}^k)$ using the previous algorithm.
 - 2.c Set $\mathbf{x}^{k+1} = \mathbf{x}^k + \alpha_k \mathbf{p}^k$, where α_k satisfies the Wolfe conditions.
if $k > m$, discard the pair $\{\mathbf{s}^{k-m}, \mathbf{p}^{k-m}\}$ from storage.
 - 2.d Compute and store $\mathbf{s}^k = \mathbf{x}^{k+1} - \mathbf{x}^k$, $\mathbf{y}^k = \nabla f(\mathbf{x}^{k+1}) - \nabla f(\mathbf{x}^k)$.

Warning

L-BFGS updates does not guarantee positive semidefiniteness of the variable metric \mathbf{H}_k in contrast to BFGS.

Example: Logistic regression - numerical results



Parameters

- ▶ For BFGS, L-BFGS and Newton's method: maximum number of iterations 200, tolerance 10^{-6} . L-BFGS memory $m = 50$.
- ▶ For accelerated gradient method: maximum number of iterations 20000, tolerance 10^{-6} .
- ▶ Ground truth: Get a high accuracy approximation of x^* and f^* by applying Newton's method for 200 iterations.

Performance of optimization algorithms

Time-to-reach ϵ

time-to-reach ϵ = number of iterations to reach ϵ \times per iteration time

The **speed** of numerical solutions depends on two factors:

- **Convergence rate** determines the number of iterations needed to obtain an ϵ -optimal solution.
- **Per-iteration time** depends on the information oracles, implementation, and the computational platform.

In general, convergence rate and per-iteration time are inversely proportional.

Finding the **fastest** algorithm is tricky! A non-exhaustive illustration:

Assumptions on f	Algorithm	Convergence rate	Iteration complexity
Lipschitz-gradient $f \in \mathcal{F}_{L,1}^{2,1}(\mathbb{R}^p)$	Gradient descent	Sublinear ($1/k$)	One gradient
	Accelerated GD	Sublinear ($1/k^2$)	One gradient
	Quasi-Newton	Superlinear	One gradient, rank-2 update
	Newton method	Sublinear ($1/k$), Quadratic	One gradient, one linear system
Strongly convex, smooth $f \in \mathcal{F}_{L,\mu}^{2,1}(\mathbb{R}^p)$	Gradient descent	Linear (e^{-k})	One gradient
	Accelerated GD	Linear (e^{-k})	One gradient
	Quasi-Newton	Superlinear	One gradient, rank-2 update
	Newton method	Linear (e^{-k}), Quadratic	One gradient, one linear system

Performance of optimization algorithms

A non-exhaustive comparison:

Assumptions on f	Algorithm	Convergence rate	Iteration complexity
Lipschitz-gradient $f \in \mathcal{F}_{L,1}^{2,1}(\mathbb{R}^p)$	Gradient descent	Sublinear ($1/k$)	One gradient
	Accelerated GD	Sublinear ($1/k^2$)	One gradient
	Quasi-Newton Newton method	Superlinear Sublinear ($1/k$), Quadratic	One gradient, rank-2 update One gradient, one linear system
Strongly convex, smooth $f \in \mathcal{F}_{L,\mu}^{2,1}(\mathbb{R}^p)$	Gradient descent	Linear (e^{-k})	One gradient
	Accelerated GD	Linear (e^{-k})	One gradient
	Quasi-Newton Newton method	Superlinear Linear (e^{-k}), Quadratic	One gradient, rank-2 update One gradient, one linear system

Accelerated gradient descent:

$$\begin{aligned}\mathbf{x}^{k+1} &= \mathbf{y}^k - \alpha \nabla f(\mathbf{y}^k) \\ \mathbf{y}^{k+1} &= \mathbf{x}^{k+1} + \gamma_{k+1}(\mathbf{x}^{k+1} - \mathbf{x}^k).\end{aligned}$$

for some proper choice of α and γ_{k+1} .

Performance of optimization algorithms

A non-exhaustive comparison:

Assumptions on f	Algorithm	Convergence rate	Iteration complexity
Lipschitz-gradient $f \in \mathcal{F}_L^{2,1}(\mathbb{R}^p)$	Gradient descent	Sublinear ($1/k$)	One gradient
	Accelerated GD	Sublinear ($1/k^2$)	One gradient
	Quasi-Newton	Superlinear	One gradient, rank-2 update
Strongly convex, smooth $f \in \mathcal{F}_{L,\mu}^{2,1}(\mathbb{R}^p)$	Newton method	Sublinear ($1/k$), Quadratic	One gradient, one linear system
	Gradient descent	Linear (e^{-k})	One gradient
	Accelerated GD	Linear (e^{-k})	One gradient
	Quasi-Newton	Superlinear	One gradient, rank-2 update
	Newton method	Linear (e^{-k}), Quadratic	One gradient, one linear system

Main computations of the Quasi-Newton method, which we will discuss in the sequel

$$\mathbf{p}^k = -\mathbf{B}_k^{-1} \nabla f(\mathbf{x}^k),$$

where \mathbf{B}_k^{-1} is updated at each iteration by adding a rank-2 matrix.

Performance of optimization algorithms

A non-exhaustive comparison:

Assumptions on f	Algorithm	Convergence rate	Iteration complexity
Lipschitz-gradient $f \in \mathcal{F}_L^{2,1}(\mathbb{R}^p)$	Gradient descent	Sublinear ($1/k$)	One gradient
	Accelerated GD	Sublinear ($1/k^2$)	One gradient
	Quasi-Newton	Superlinear	One gradient, rank-2 update
	Newton method	Sublinear ($1/k$), Quadratic	One gradient, one linear system
Strongly convex, smooth $f \in \mathcal{F}_{L,\mu}^{2,1}(\mathbb{R}^p)$	Gradient descent	Linear (e^{-k})	One gradient
	Accelerated GD	Linear (e^{-k})	One gradient
	Quasi-Newton	Superlinear	One gradient, rank-2 update
	Newton method	Linear (e^{-k}), Quadratic	One gradient, one linear system

The main computation of the Newton method requires the solution of the linear system

$$\nabla^2 f(\mathbf{x}^k) \mathbf{p}^k = -\nabla f(\mathbf{x}^k) .$$

References |

- [1] JE Dennis and Jorge J Moré.
A characterization of superlinear convergence and its application to quasi-newton methods.
Mathematics of Computation, 28(126):549–560, 1974.
- [2] J. Nocedal and S.J. Wright.
Numerical Optimization.
Springer, 2006.